

Distributed Computing: Autumn 2018
Programming Assignment 1: Implementing Vector Clocks
Submission Date: ~~23rd August 2018~~ 25th August, 9:00 pm

Goal: The goal of this assignment is to implement vector-clocks and Singhal-Kshemkalyani optimization on a Distributed System. Implement both these locking algorithms in C++. Then, you have to compare the overheads incurred with message stored and exchanged.

Details. You are given as input a system of n nodes (processes) connected to each other in the form of a graph topology. These nodes communicate with their neighbors (in the graph) through messages. You can implement these nodes as processes/threads communicating through sockets.

As studied in the class, each process executes three events: internal, message send & message receive. To simulate these events you can assume the following: each process creates and executes internal and message send events with a delay that is exponentially distributed with inter-event time λ ms. Assume that the ratio of internal to message send events on each process is α (which for instance can be 1.5). The application terminates when each processes has sent a total of m messages.

In this setting, implement the normal vector-clocks and Singhal-Kshemkalyani's optimization of vector-clocks. Then demonstrate the savings in message communication obtained by using the optimization.

Input: The input to the program will be a file, named inp-params.txt, consisting of all the parameters described above and the graph topology. The first line of the input file will contain the parameters $:n, \lambda, \alpha, m$. For example, $n = 15, \lambda = 5, \alpha = 1.5, m = 40$

The second line onwards, will contain the graph topology in the form of an adjacency list. For instance if n is 3, a sample topology showing a complete graph is as follows:

```
1 2 3
2 1 3
3 1 2
```

Output: Your program output should demonstrate the strong consistency property of vector-clocks: If two events x and y have timestamps vh and vk , respectively, then

$$x \rightarrow y \Leftrightarrow vh < vk \quad (1)$$

$$x \parallel y \Leftrightarrow vh \parallel vk \quad (2)$$

To demonstrate this, you have to output the contents of all the events onto a common logfile with two time-stamps: real time-stamps and vector time-stamps. Since you will be executing all these programs on the same machine or in a LAN environment like IITH, you can assume that the clocks of the machines are synchronized.

The contents of the logfile should be as follows for both the algorithms:

Process1 executes internal event e11 at 10:00, vc: [1 0 0 0]

Process2 executes internal event e21 at 10:01, vc: [0 1 0 0]

Process3 sends message m31 to process2 at 10:02, vc: [0 0 1 0]

Process3 executes internal event e32 at 10:03, vc: [0 0 2 0]

Process2 executes internal event e22 at 10:04, vc: [0 2 0 0]

Process2 receives m31 from process3 at 10:05, vc: [0 3 1 0]

.
.
.

The output should demonstrate Eqn(1) and Eqn(2). **In addition to this, you have to output space utilized by each process for storing the vector clocks and sending messages by both the algorithms.**

Report: You have to submit a report for this assignment. As mentioned earlier, this report should contain a comparison of the performance of vector-clocks and Singhal-Kshemkalyani optimization. You must run both these algorithms multiple times to compare the performances and display the result in form of a graph.

You run both these algorithms varying the number of processes from 10 to 15 in the increments of 1 while keeping other parameters same. You can assume any topology for creating the graph with each of these processes. Assume that m is 50 for all the executions.

You measure the average number of vector-clock entries sent in each message. Clearly this number is going to be fixed for the normal vector-clock. But for Singhal-Kshemkalyani optimization, this number should be lesser.

The graph in the report will be as follows: the x-axis will vary the number of threads from 10 to 15 in the increments of 1 (as explained above) while the y-axis will show the average number of entries in each message sent. Finally, you must also give an analysis of the results while explaining any anomalies observed.

Deliverables: You have to submit the following:

- The source files containing the actual program to execute. The normal vector clock implementation should be named as VC-<rollno>.cpp and Singhal Kshemkalyani's optimization as SK-<rollno>.cpp.
- A readme.txt that explains how to execute the program.
- The report as explained above.

Zip all the files and name it as ProgAssn1-<rollno>.zip. Please follow the naming convention strictly. Otherwise, your submission will not be evaluated. Then upload the zip on the google classroom page of this course. Submit it by ~~23rd August 2018~~ **25th August, 9:00 pm.**